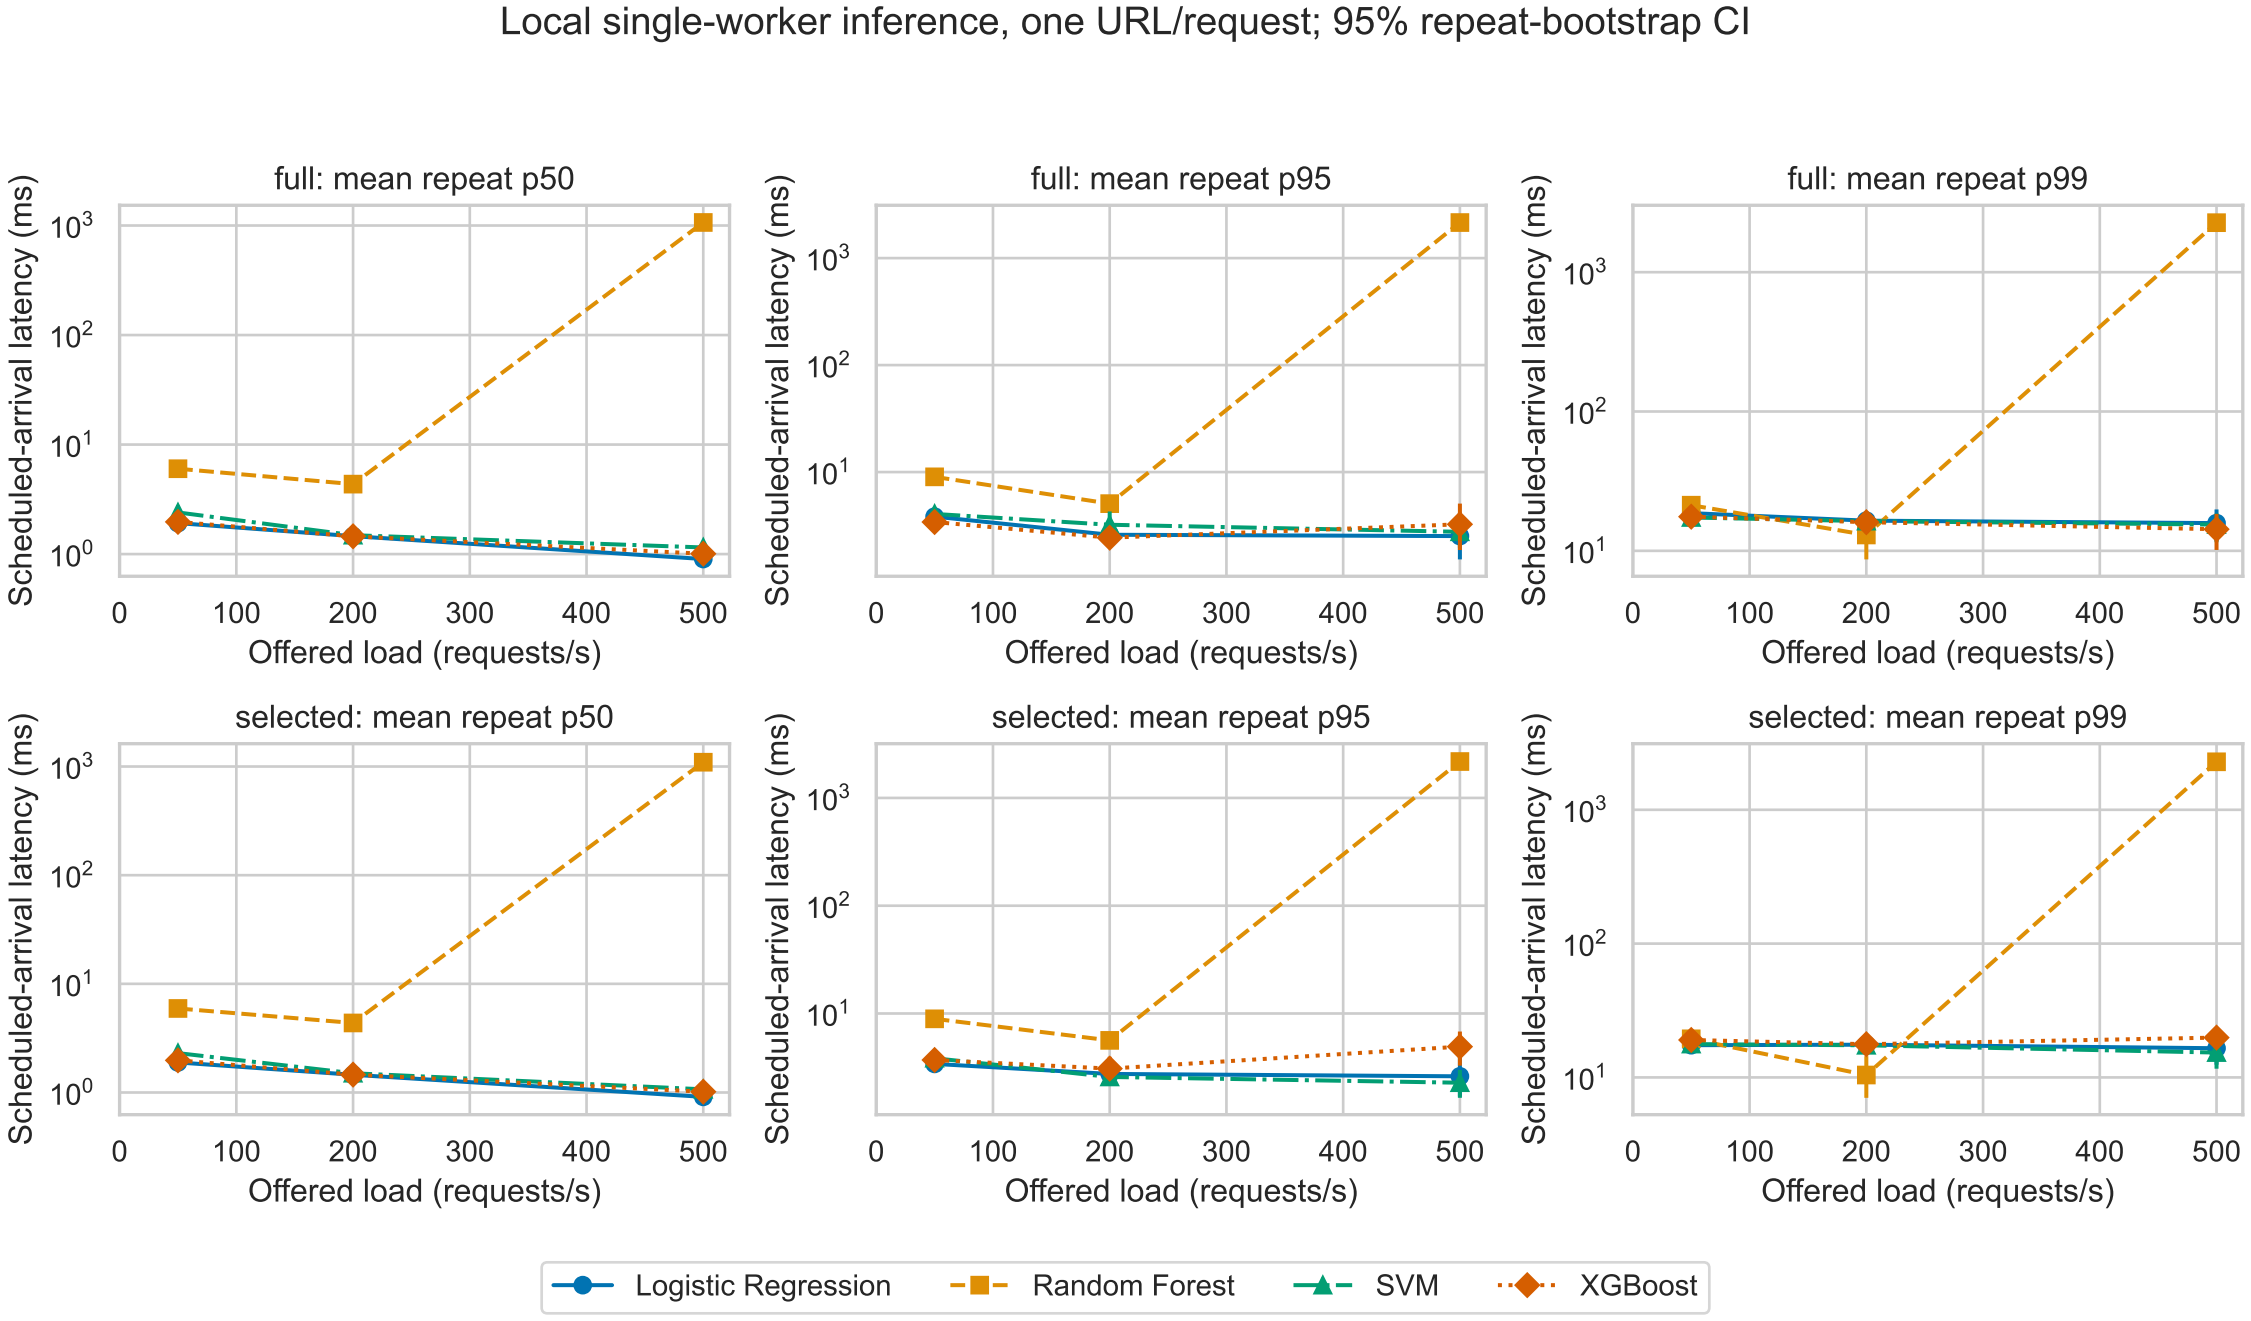
Local single-worker inference, one URL/request; 95% repeat-bootstrap CI
full: mean repeat p50
full: mean repeat p95
full: mean repeat p99
selected: mean repeat p50
selected: mean repeat p95
selected: mean repeat p99
Scheduled-arrival latency (ms)
Offered load (requests/s)
Logistic Regression
Random Forest
SVM
XGBoost